

Soumya Gupta

soumya0343@gmail.com | [+91-8208897490](tel:+91-8208897490) | linkedin.com/in/soumya-gupta | github.com/soumya0343 | [Portfolio](#)

EDUCATION

B.E. in Electronics and Communications Engineering

Birla Institute Of Technology and Science, Pilani, K. K. Birla Goa Campus

Oct 2022 – May 2026

Goa, India

WORK EXPERIENCE

Software Development Engineer

Jul 2026 - Present

Dezerv Investments

Bengaluru, India

- Revamping the onboarding flow in the Dezerv app, architected as a **plug-and-play module** so any surface needing Dezerv onboarding reuses one flow instead of reimplementing it; on release it serves **100% of Dezerv clients**.
- Building backend and frontend **AI pipeline** serving foundations in **Python** that the wider engineering team builds on, alongside **Go**, **React**, and **Flutter/Dart**.

Full Stack Developer Intern

Dec 2025 - Feb 2026

Chakra Tech

Remote

- Developed an **MCP server** (~20 tools) powering an AI assistant for contextual market queries, designing the tool schemas and tool-calling surface over live market data.
- Engineered an **event-driven ingestion pipeline** consuming Polymarket **WebSocket** streams; used **BullMQ on Redis** for async job orchestration with retries, deduplication, and dead-letter handling.
- Architected a **hybrid storage layer**, **ClickHouse** for high-volume trade event ingestion (10k-row batches, partitioned by month) and **PostgreSQL** for transactional metadata.
- Resolved **N+1 query bottlenecks** via query-plan analysis and targeted indexing; reduced endpoint latency from **~15s to <100ms**.

Engineering Intern

Jul 2025 - Dec 2025

Dezerv Investments

Bengaluru, India

- Architected a scalable **Go** backend to build a self-serve, automated platform for SEBI-compliant reporting for **5,000+** clients using **AWS S3** and **Argo Workflows**, cutting manual reporting time by **~83%**.
- Shipped a reusable **React** operations dashboard introducing direct-to-S3 bulk uploads, driving **~70%** more reports per cycle.
- Collaborated cross-functionally in an **Agile** environment with product managers and analysts to roll out a **React** Rationale Dashboard, reducing client queries by **~35%**.
- Contributed to Dezerv's first Flutter Web client dashboard (PMS), enabling desktop access for **16.9% non-app users**.

Tech Intern

May 2025 - Jun 2025

Jobslet

Remote

- Directed development of analytics dashboards using **React** and **Node.js** to track stage-wise user progression and first-time vs. repeat activity, allowing ops teams to pinpoint recruitment bottlenecks and cutting manual reporting effort by **~40%**.
- Optimized backend data pipelines to cut latency by **~70%**, enabling near real-time updates for recruitment workflows.
- Executed **Elasticsearch**-backed indexing, filtering, and optimized search logic, reducing recruiter shortlisting time by **35%**.
- Improved **LLM-based** resume-parsing pipelines for better field extraction and scoring-based candidate-job matching.

PROJECTS

Curate - Intent-Driven Product Recommendation Engine | [Live](#) | [GitHub](#)

- Designed a **5-stage hybrid pipeline** bracketing two LLM calls (intent, rerank) around deterministic filter/retrieve/prerank, decomposing a query into **sub-needs** with per-need **vector search** so distinct intents never collapse into one averaged embedding.
- Built **provenance trust tiers** over an Amazon India catalogue (**1.59M** rows), gating title-verified attributes as fact while inferred ones only nudge ranking, with assumption chips and relaxation notices surfaced to users.
- Shipped **FastAPI**/Pydantic v2 with **SSE-over-POST** streaming to a **React 19**/TS client over an in-memory **NumPy** catalogue, plus a **multi-provider** chain (Gemini/Groq/Cerebras) with key rotation on 429; **270 tests**.

Saral - Multilingual Multi-Agent Support-Resolution System | [GitHub](#)

- Authored a **LangGraph** supervisor-orchestrated **multi-agent** graph from scratch (triage, RAG, action, synthesis) with parallel mixed-intent fan-out, tool calling, and **suspend/resume** conversation memory.
- Implemented per-customer **RAG** citing the user's policy clause, over a **Sarvam**/Anthropic provider abstraction, in **EN/Hindi/Hinglish**.
- Devised an **eval pipeline**, **190** labeled scenarios with a per-language LLM judge scoring groundedness, cross-lingual consistency, impersonation, and injection resistance.
- Orchestrated an **event-driven** Redis Streams runtime with checkpoint/resume, **exactly-once** writes, token identity, and step-up OTP.

RepoLens - SDLC Intelligence Platform | [GitHub](#)

- Deployed a live LLM assistant exposing **6 tool-use tools** over repo analytics, with **multi-provider** inference (Groq/Gemini), Redis caching, and structured-output fallback on failure.
- Structured an **event-driven** backend, FastAPI and 5 specialized **ARQ/Redis** workers, over async **PostgreSQL/TimescaleDB** and a **Neo4j** collaboration graph.
- Unified 6 analysis engines into UnifiedRiskScorer, FP-Growth coupling, Tree-sitter AST violations, bus-factor, and **Neo4j** reviewer graph.
- Shipped **45 REST endpoints** across **6** services to **Azure** (Docker, Caddy auto-TLS, **GitHub Actions CI**, OAuth/JWT).

TECHNICAL SKILLS

Programming Languages: Python, TypeScript, JavaScript, Go, C/C++, Java, Dart, SQL

AI & LLM: LangGraph, Multi-Agent Systems, RAG, Tool Calling, MCP Server, Prompt Engineering, LLM Eval, Embeddings, SSE

Frameworks & Technologies: FastAPI, Node.js, Express, React, Next.js, Flutter, Pydantic, HTML/CSS

Databases: PostgreSQL, TimescaleDB, MongoDB, Redis, Neo4j, ClickHouse, Elasticsearch, Firebase

Tools & Platforms: Git/GitHub, Docker, Kubernetes, Azure, AWS, GitHub Actions, Argo Workflows, ARQ, BullMQ, Caddy, Grafana

Concepts: System Design, Microservices, REST APIs, WebSockets, OAuth/JWT, Data Structures & Algorithms, OOP, SDLC