

Soumya Gupta

soumya0343@gmail.com | [+91-8208897490](tel:+91-8208897490) | [LinkedIn](#) | [GitHub](#) | [Portfolio](#)

EDUCATION

B.E. in Electronics and Communications Engineering

Birla Institute Of Technology and Science, Pilani, K. K. Birla Goa Campus

Oct 2022 – May 2026

Goa, India

WORK EXPERIENCE

Full Stack Developer Intern

Dec 2025 - Feb 2026

Chakra Tech

Remote

- Developed an **MCP server** (~20 tools) powering an AI assistant for contextual market queries, designing the tool schemas and tool-calling surface over live market data.
- Built an **event-driven ingestion pipeline** consuming Polymarket **WebSocket** streams; used **BullMQ on Redis** for async job orchestration with retries, deduplication, and dead-letter handling.
- Architected a **hybrid storage layer**, **ClickHouse** for high-volume trade event ingestion (10k-row batches, partitioned by month) and **PostgreSQL** for transactional metadata.
- Resolved **N+1 query bottlenecks** via query-plan analysis and targeted indexing; reduced endpoint latency from ~15s to <100ms.

Engineering Intern

Jul 2025 - Dec 2025

Dezerv Investments

Bengaluru, India

- Architected a scalable **Go** backend to build a self-serve, automated platform for SEBI-compliant reporting for **5,000+** clients using **AWS S3** and **Argo Workflows**, cutting manual reporting time by ~83%.
- Shipped a reusable **ReactJs** operations dashboard introducing direct-to-S3 bulk uploads, driving ~70% more reports per cycle.
- Collaborated cross-functionally in an **Agile** environment with product managers and analysts to roll out a **ReactJs** Rationale Dashboard, reducing client queries by ~35%.
- Contributed to Dezerv's first Flutter Web client dashboard (PMS), enabling desktop access for **16.9% non-app users**.

Tech Intern

May 2025 - Jun 2025

Jobslet

Remote

- Directed development of analytics dashboards using **ReactJs** and **Node.js** to track stage-wise user progression and first-time vs. repeat activity, allowing ops teams to pinpoint recruitment bottlenecks and cutting manual reporting effort by ~40%.
- Optimized backend data pipelines to cut latency by ~70%, enabling near real-time updates for recruitment workflows.
- Executed **Elasticsearch**-backed indexing, filtering, and optimized search logic, reducing recruiter shortlisting time by 35%.
- Improved **LLM-based** resume parsing pipelines for better field extraction and scoring-based candidate-job matching.

Software Developer Intern

May 2024 - Jul 2024

Multigraphics Group

Delhi, India

- Launched a Laravel-based Employee Management System with CRUD APIs and authentication to centralize internal operations for **50+** concurrent users.
- Produced a secure, responsive UI with advanced filtering to eliminate slow data access, cutting latency by ~40%.
- Developed a Question Paper Management System supporting creation, editing, and formatting workflows for **50K+** documents annually.
- Added logging and audit trails with secure PDF exports to address missing traceability, improving accountability by ~25%.

PROJECTS

Saral - Multilingual Multi-Agent Support-Resolution System | [GitHub](#)

- Authored a **LangGraph** supervisor-orchestrated **multi-agent** graph from scratch (triage, RAG, action, synthesis) with parallel mixed-intent fan-out, tool calling, and **suspend/resume** conversation memory.
- Built per-customer **RAG** citing the user's policy clause, over a **Sarvam**/Anthropic provider abstraction, in **EN/Hindi/Hinglish**.
- Designed an **eval pipeline**, **190** labeled scenarios with a per-language LLM judge scoring groundedness, cross-lingual consistency, impersonation, and injection resistance.
- Architected an **event-driven** Redis Streams runtime with checkpoint/resume, **exactly-once** writes, token identity, and step-up OTP.

InferLog - LLM Inference Observability Platform | [GitHub](#)

- Deployed a **multi-provider** LLM serving layer (Groq/Gemini/OpenAI/Anthropic/**Sarvam**) via a pluggable **BaseProvider** SDK with **SSE streaming** and TTFT tracking.
- Decoupled logging from the hot path with **Redis Streams** consumer groups and **XAUTOCLAIM** crash recovery.
- Two-stage **Presidio PII redaction** with p50/p95/p99 analytics via a concurrently-refreshed **materialized view**.
- Containerized with **19 Kubernetes** manifests (HPA, consumer-group scaling) and **GitHub Actions CI** publishing images to GHCR.

RepoLens - SDLC Intelligence Platform | [GitHub](#)

- Deployed a live LLM assistant exposing **6 tool-use tools** over repo analytics, with **multi-provider** inference (Groq/Gemini), Redis caching, and structured-output fallback on failure.
- Architected an **event-driven** backend, FastAPI and 5 specialized **ARQ/Redis** workers, over async **PostgreSQL/TimescaleDB** and a **Neo4j** collaboration graph.
- Built 6 analysis engines into UnifiedRiskScorer, FP-Growth coupling, Tree-sitter AST violations, bus-factor, and **Neo4j** reviewer graph.
- Shipped to production on **Azure** (Docker, Caddy auto-TLS, **GitHub Actions CI**) with OAuth/JWT.

TECHNICAL SKILLS

Programming Languages: Python, TypeScript, JavaScript, Go/GoLang, C/C++, Java, Dart, SQL

AI & LLM: LangGraph, Multi-Agent Systems, RAG, Tool Calling, MCP Server, Prompt Engineering, LLM Eval, Embeddings, SSE

Frameworks & Technologies: FastAPI, NodeJs, Express, ReactJs, NextJs, Flutter, Pydantic, HTML/CSS

Databases: PostgreSQL, TimescaleDB, MongoDB, Redis, Neo4j, ClickHouse, Elasticsearch, Firebase

Tools & Platforms: Git/GitHub, Docker, Kubernetes, Azure, AWS, GitHub Actions, Argo Workflows, ARQ, BullMQ, Caddy, Grafana

Concepts: System Design, Microservices, REST APIs, WebSockets, OAuth/JWT, Data Structures & Algorithms, OOP, SDLC